

B04 Targeted Remediation Cards

Use only the card matching the failed competency. Save the failed evidence first. Repair does not auto-open Gate B or C.

Card	Weak competency	Fresh repair task	Pass evidence
R1	Workload/source boundary	Reclassify four fresh activities as OLTP/OLAP and state one source-protection rule.	Justify from access pattern + business risk.
R2	Storage fit	Choose warehouse/lake/lakehouse for three fresh cases and reject one alternative.	Cite shape, consumer, governance, access + replay.
R3	History/key strategy	Repair a duplicate effective-date history and distinguish business vs surrogate key.	No ambiguous active version; history rule explicit.
R4	ETL/ELT/layer boundary	Draw a fresh source -> landing -> transform -> serve flow.	Raw evidence + transformation boundary unambiguous.
R5	Latency decision	Choose batch/micro-batch/streaming for three SLAs.	Late data + replay addressed; faster not assumed better.
R6	Lineage/ownership	Trace one wrong KPI upstream and list minimum run metadata.	Owner, run ID, freshness + downstream impact visible.
R7	Quality severity	Classify five defects as BLOCK/WARN/QUARANTINE with owner/action.	No silent deletion + independent reconciliation.
R8	Recovery/replay	Write a bounded backfill after a partial run.	State protected + post-recovery proof independent.
R9	Integrated defense	Defend one architecture in 90 seconds from source to serving and failure recovery.	Every major choice has evidence/trade-off, not tool-name logic.

Reassessment rule

A targeted repair proves only the named weak competency. If a fresh full Gate is required, it must be explicitly assigned. Gate B/C remain protected, distinct full reassessments; passing a micro-repair never substitutes for them.

Append-only evidence

Keep the original failed attempt, failure code, repair work and fresh proof. Do not overwrite history when the later result passes.